

IT308: Operating Systems

Virtual memory: Paging, TLB

How to fix both internal and external fragmentation?



Today: small fixed-size luggage



Paging

- Instead of dealing with variable sized segments, we should use fixed size chunks
- These are called pages

- Instead of dealing with variable sized segments, we should use fixed size chunks
- These are called pages
- We divide both virtual memory and physical memory into pages of equal size (typically 4kB)

- Instead of dealing with variable sized segments, we should use fixed size chunks
- These are called pages
- We divide both virtual memory and physical memory into pages of equal size (typically 4kB)
- Physical memory pages are called page frames

Example

0	reserved for OS	page frame 0 of physical memory
16	(unused)	page frame 1
32	page 3 of AS	page frame 2
48	page 0 of AS	page frame 3
64	(unused)	page frame 4
80	page 2 of AS	page frame 5
96	(unused)	page frame 6
112	page 1 of AS	page frame 7
128		

Free Space Management

- The OS can keep a list of free pages
- Allocation of a page is $O(1)$

Mapping Virtual to Physical

- We need to keep a map of virtual pages to physical pages
- This is called a page table and there is one per process

Mapping Virtual to Physical

- We need to keep a map of virtual pages to physical pages
- This is called a page table and there is one per process
- Let's start with a simple example with 16 byte pages and 6 bit addresses

Va5	Va4	Va3	Va2	Va1	Va0
-----	-----	-----	-----	-----	-----

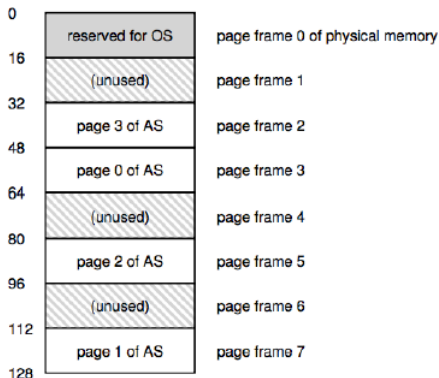

Question

If we have 6 bit virtual addresses and 16 byte pages, how many bits of the virtual address tells us which virtual page an address is?

- (A) 0
- (B) 1
- (C) 2
- (D) 3
- (E) 4


Example

- `movl 21, %eax`
- $21 = 010101$
- So the 5th byte of the virtual page 1
- Which is the 5th byte of physical page 7



Summary of Address Translation

- The physical page is looked up in the page table at index virtual page number
- Then at the offset within that physical page



Question

If we have a 32 bit machine with 4kB pages and each page table entry is 4 bytes, how big is each process page table?

- (A) 32 bits
- (B) 2^{20} bytes
- (C) 2^{22} bytes
- (D) 2^{32} bytes

Question

If we have a 32 bit machine with 4kB pages and each page table entry is 4 bytes, how big is each process page table?


- (A) 32 bits
- (B) 2^{20} bytes
- (C) 2^{22} bytes
- (D) 2^{32} bytes

So 2^{22} bytes and 100 processes is 400 MB

Where do we store page tables?

- In physical memory (as opposed to in a CPU register)
- A simple, linear Page Table is just an array
- Each entry in the array is a Page Table Entry (PTE)

What are we keeping in the PTE?



- Page Frame Number (PFN), the most important part!
- Protection bits (r/w), Present bit (swapped out), Referenced bit (recently accessed)
 - More on these later

How long does it take?

- On each load or store from/to memory, we have to fetch a PTE from memory
- So each memory access is now two memory access?
- So page tables are **big and slow** . . .

Enter: the TLB!

- TLB is a cache of the page table
- Each TLB entry should cache all the information in a PTE
- It also needs to store the VPN as a tag
 - To be used when the hardware searches TLB for a particular VPN
- A typical TLB entry

VPN | PFN | other bits

Reading an Array (no TLB)

- Suppose we have a 10KB array of integers
 - Assume 4KB pages
- With no TLB, how many memory accesses are required to read the whole array?


Reading an Array (no TLB)

- Suppose we have a 10KB array of integers
 - Assume 4KB pages
- With no TLB, how many memory accesses are required to read the whole array?
 - $10\text{KB} / 4 = 2560$ integers in the array
 - Each requires one page table lookup, one memory read
 - 5120 reads, plus more for the instructions themselves

Reading an Array (with TLB)

- Same example, now with TLB
 - 10KB integer array
 - 4KB pages
 - Assume the TLB starts off cold (i.e. empty)
- How many memory accesses to read the array?
 - 2560 to read the integers
 - 3 page table lookups
 - 2563 total reads

Process 1's View of Virtual Memory



TLB

VPN	PFN

- TLB, like any cache, is effective because of locality

- TLB, like any cache, is effective because of locality
- Spatial locality: if you access memory address x , it is likely you will access $x + 1$ soon
 - Most of the time, x and $x + 1$ are in the same page

- TLB, like any cache, is effective because of locality
- Spatial locality: if you access memory address x , it is likely you will access $x + 1$ soon
 - Most of the time, x and $x + 1$ are in the same page
- Temporal locality: if you access memory address x , it is likely you will access x again soon
 - The page containing x will still be in the TLB, hopefully