

# IT308: Operating Systems

## Implementing Locks

# Locks and Scheduling

- Threads as entities scheduled by OS
- Locks provide some minimal amount of control over scheduling (back) to programmers
  - e.g., guarantee that no more than a single thread can ever be active within that code (critical section)
- In a previous course, we studied how to use Pthread locks

# Locks and Scheduling

- Threads as entities scheduled by OS
- Locks provide some minimal amount of control over scheduling (back) to programmers
  - e.g., guarantee that no more than a single thread can ever be active within that code (critical section)
- In a previous course, we studied how to use Pthread locks
  - But how should we build a lock?

# Building a Lock

- Over the years, different hardware primitives have been added to the instruction sets of various computer architectures
- Will study how to use them in order to build a lock
- Will study how OS gets involved to complete the picture and enable us to build a sophisticated locking library

# Requirements for Locks

- Correctness
  - Only one thread in critical section at a time

# Requirements for Locks

- Correctness
  - Only one thread in critical section at a time
- Fairness
  - Each thread gets a fair chance at acquiring the lock

# Requirements for Locks

- Correctness
  - Only one thread in critical section at a time
- Fairness
  - Each thread gets a fair chance at acquiring the lock
- Performance
  - Time overhead for a lock without and with contentions (possibly on multiple CPUs)

# Mutual Exclusion and Interrupts

- Mutual exclusion
  - Want to prevent another thread from executing while current thread is in critical section
- Solution: disable interrupts for critical sections



# Mutual Exclusion and Interrupts

- Mutual exclusion
  - Want to prevent another thread from executing while current thread is in critical section
- Solution: disable interrupts for critical sections
  - before entering critical section, disable interrupts
  - after exiting critical section, reenale interrupts
  - This enforces mutual exclusion

# Lock Implementation 1

```
void lock ()
{
    DisableInterrupt(); // special hardware instruction
}

void unlock()
{
    EnableInterrupt();  // special hardware instruction
}
```

- Negatives?

# Lock Implementation 1

```
void lock ()
{
    DisableInterrupt(); // special hardware instruction
}

void unlock()
{
    EnableInterrupt();  // special hardware instruction
}
```

- Negatives?
  - allow calling thread to perform privileged operations – “trust?”

# Lock Implementation 1

```
void lock ()
{
    DisableInterrupt(); // special hardware instruction
}

void unlock()
{
    EnableInterrupt();  // special hardware instruction
}
```

- Negatives?
  - allow calling thread to perform privileged operations – “trust?”
  - disables multiprogramming even if another thread is NOT interested in critical section

## Lock Implementation 2

```
typedef struct __lock_t { int flag; } lock_t;

void init(lock_t *mutex) {
    mutex->flag = 0; // 0 -> lock is available, 1 -> held
}

void lock(lock_t *mutex) {
    while (mutex->flag == 1) // TEST the flag
        ; // spin-wait (do nothing)
    mutex->flag = 1;          // now SET it!
}

void unlock(lock_t *mutex) {
    mutex->flag = 0;
}
```

What problems does this solution have?

# Malicious Scheduler

- Pretend being a malicious scheduler, one that interrupts threads at the most inopportune of times in order to foil synchronization primitives
- Although the exact sequence of interrupts may be improbable, it is possible, and that is all we need to demonstrate that a particular approach does not work

# Trace: No Mutual Exclusion

## Thread 1

call lock()

while (flag == 1) **false**

**interrupt: switch to Thread 2**

flag = 1; // set flag to 1 (too!)

## Thread 2

call lock()

while (flag == 1) **false**

flag = 1;

**interrupt: switch to Thread 1**

# Atomic Test-And-Set

- machine instruction used to synchronize threads
  - atomically sets the value of a specified memory location and places that memory location's old value into a register
- A C description of the machine instruction

```
int TestAndSet(int *old_ptr, int new) {  
    int old = *old_ptr; // fetch old value at old_ptr  
    *old_ptr = new;      // store 'new' into old_ptr  
    return old;          // return the old value  
}
```

- ...all this is done atomically



# Lock Implementation 3

```
typedef struct __lock_t { int flag; } lock_t;

void init(lock_t *lock) {
    lock->flag = 0; // 0 -> lock is available, 1 -> held
}

void lock(lock_t *lock) {
    while (TestAndSet(&lock->flag, 1) == 1) // TEST the flag
        ; // spin-wait (do nothing)
}

void unlock(lock_t *lock) {
    lock->flag = 0;
}
```

- As long as the lock is held by another thread, TestAndSet() will repeatedly return 1, and thus the calling thread will spin-wait

# Drawbacks of Spin Lock

- Starvation is possible
  - a waiting thread may wait (spin) forever under contention
- Spinning can be costly
  - assume thread holding lock is preempted within a critical section
  - subsequently if  $N-1$  waiting threads are scheduled by the OS, then each of those will waste CPU cycles for the duration of a time slice before giving up