

IT308: Operating Systems

Fine-grained locking

Recall: Semaphores

- $P()$ or $\text{wait}()$ or $\text{down}()$
 - From Dutch “proeberen”, meaning “test”
 - Atomic action:
 - Wait for semaphore value to become > 0 , then decrement it
- $V()$ or $\text{signal}()$ or $\text{up}()$
 - From Dutch “verhogen”, meaning “increment”
 - Atomic action:
 - Increments semaphore value by 1.

Simple Semaphore Implementation

```
struct semaphore {  
    int val;  
    thread_list waiting; // List of threads waiting for semaphore  
}
```

```
P(semaphore Sem):    // Wait until > 0 then decrement  
    while (Sem.val <= 0) {  
        add this thread to Sem.waiting;  
        block(this thread); // What does this do??  
    }  
    Sem.val = Sem.val - 1;  
    return;
```

```
V(semaphore Sem):    // Increment value and wake up next thread  
    Sem.val = Sem.val + 1;  
    if (Sem.waiting is nonempty) {  
        remove a thread T from Sem.waiting;  
        wakeup(T);  
    }
```

P() and V() must be atomic actions!

Simple Semaphore Implementation

```
struct semaphore {  
    int val;  
    thread_list waiting; // List of threads waiting for semaphore  
}
```

```
P(semaphore Sem): // Wait until > 0 then decrement  
    while (Sem.val <= 0) {  
        add this thread to Sem.waiting;  
        block(this thread); // What does this do??  
    }  
    Sem.val = Sem.val - 1;  
    return;
```

← Why is this a while loop and not just an if statement?

```
V(semaphore Sem): // Increment value and wake up next thread  
    Sem.val = Sem.val + 1;  
    if (Sem.waiting is nonempty) {  
        remove a thread T from Sem.waiting;  
        wakeup(T);  
    }
```

- Single shared object
- Want to allow any number of threads to read simultaneously
- But, only one thread should be able to write to the object at a time
 - And, not interfere with any readers ...

Readers/Writers

- readers share:
 - semaphore wrt; // initialized to 1
 - int readcount; // initialized to 0
- writers also share semaphore wrt

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

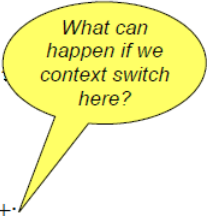
- Seems simple, but this code is broken ...

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```



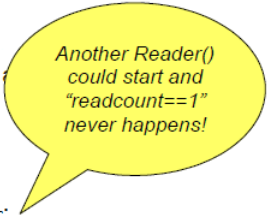
*What can
happen if we
context switch
here?*

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```



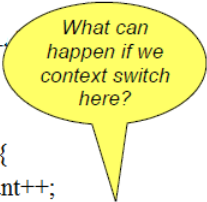
*Another Reader()
could start and
"readcount==1"
never happens!*

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```



*What can
happen if we
context switch
here?*

A Writer() could start, P the semaphore first, then subsequent Reader() threads would be able to get past the semaphore (since "readcount != 1")

Writer.

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader.

```
while (1) {  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    // reading  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
}
```

Readers/Writers fixed

- Problem: Multiple Readers are accessing readcount
 - Solution: Make “increment, test, P” and “decrement, test, V” both atomic - using a mutex

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```

Readers/Writers fixed

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```

What if a Writer() is active, the first Reader() stalls on P(wrt), and additional Readers() try to enter?

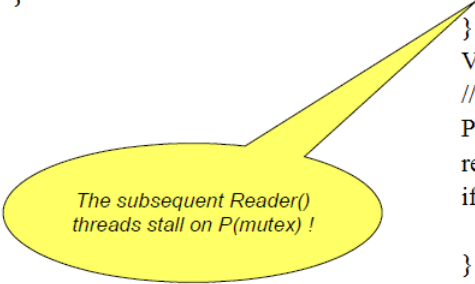
Readers/Writers fixed

Writer:

```
while (1) {  
    P(wrt);  
    // writing  
    V(wrt);  
}
```

Reader:

```
while (1) {  
    P(mutex);  
    readcount++;  
    if (readcount == 1) {  
        P(wrt);  
    }  
    V(mutex);  
    // reading  
    P(mutex);  
    readcount--;  
    if (readcount == 0) {  
        V(wrt);  
    }  
    V(mutex);  
}
```



The subsequent Reader() threads stall on P(mutex) !

A sorted linked list

```
struct Node {  
    int value;  
    Node* next;  
};
```

```
struct List {  
    Node* head;  
};
```

```
void insert(List* list, int value) {  
  
    Node* n = new Node;  
    n->value = value;  
  
    // assume case of inserting before head of  
    // of list is handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value > value)  
            break;  
  
        prev = cur;  
        cur = cur->next;  
    }  
  
    n->next = cur;  
    prev->next = n;  
}
```

What can go wrong if multiple threads operate on the linked list simultaneously?


```
void delete(List* list, int value) {  
  
    // assume case of deleting first element is  
    // handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value == value) {  
            prev->next = cur->next;  
            delete cur;  
            return;  
        }  
  
        prev = cur;  
        cur = cur->next;  
    }  
}
```

Example: concurrent insertion



- Thread 1 attempts to insert 6
- Thread 2 attempts to insert 7



Thread 1:




Example: concurrent insertion



Thread 1 and thread 2 both compute same prev and cur.
Result: one of the insertions gets lost!

Result: (assuming thread 1 updates prev->next before thread 2)



Solution 1: global locking

```
struct Node {  
    int value;  
    Node* next;  
};
```

```
struct List {  
    Node* head;  
    Lock lock;  
};
```

← Per-list lock


```
void insert(List* list, int value) {  
  
    Node* n = new Node;  
    n->value = value;  
  
    lock(list->lock);  
  
    // assume case of inserting before head of  
    // of list is handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value > value)  
            break;  
  
        prev = cur;  
        cur = cur->next;  
    }  
    n->next = cur;  
    prev->next = n;  
    unlock(list->lock);  
}
```

```
void delete(List* list, int value) {  
  
    lock(list->lock);  
  
    // assume case of deleting first element is  
    // handled here (to keep slide simple)  
  
    Node* prev = list->head;  
    Node* cur = list->head->next;  
  
    while (cur) {  
        if (cur->value == value) {  
            prev->next = cur->next;  
            delete cur;  
            unlock(list->lock);  
            return;  
        }  
  
        prev = cur;  
        cur = cur->next;  
    }  
    unlock(list->lock);  
}
```


Single global lock

- It is relatively simple to implement correct mutual exclusion for data structure operations (we just did it!)
- Disadvantages:
 - Operations on the data structure are serialized
 - May limit application performance


Hand-over-hand locking




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes




Concurrent Removes



Concurrent Removes




Concurrent Removes




Concurrent Removes


Bad news, c not removed



With Two Locks




remove(b)




remove(c)

Removing a node




27


Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node




Removing a node



Removing a node



Removing a node



Removing a node

