

# IT308: Operating Systems

## Synchronization: Deadlock characterization

# Deadlock (Definition)

- A situation in which two or more threads or processes are blocked and cannot proceed
- “blocked” either on a resource request that can’t be granted, or waiting for an event that won’t occur

## Resource deadlock: canonical example

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

# Resource deadlock: canonical example

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

One possible execution timeline:

```
lock(&A); // Thread 1 acquires lock A  
<context switch to Thread 2>
```

# Resource deadlock: canonical example

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

One possible execution timeline:

```
lock(&A); // Thread 1 acquires lock A  
<context switch to Thread 2>
```

```
lock(&B); // Thread 2 acquires lock B  
lock(&A); // Thread 2 blocks because lock A is taken  
<context switch to Thread 1>
```

# Resource deadlock: canonical example

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

One possible execution timeline:

```
lock(&A); // Thread 1 acquires lock A  
<context switch to Thread 2>
```

```
lock(&B); // Thread 2 acquires lock B  
lock(&A); // Thread 2 blocks because lock A is taken  
<context switch to Thread 1>
```

```
lock(&B); // Thread 1 blocks because lock B is taken
```

# Resource deadlock: canonical example

Thread 1:

```
lock(&A);  
lock(&B);
```

Thread 2:

```
lock(&B);  
lock(&A);
```

One possible execution timeline:

```
lock(&A); // Thread 1 acquires lock A  
<context switch to Thread 2>
```

```
lock(&B); // Thread 2 acquires lock B  
lock(&A); // Thread 2 blocks because lock A is taken  
<context switch to Thread 1>
```

```
lock(&B); // Thread 1 blocks because lock B is taken
```

Both threads are waiting for each other: they are **deadlocked**

# System Modeling

- Set of resource types:  $R_1, R_2, \dots, R_m$ 
  - There are multiple resource of each type: e.g., 3 NICs, 4 disks
- Set of processes (or threads):  $P_1, P_2, \dots, P_n$



# System Modeling

- Set of resource types:  $R_1, R_2, \dots, R_m$ 
  - There are multiple resource of each type: e.g., 3 NICs, 4 disks
- Set of processes (or threads):  $P_1, P_2, \dots, P_n$
- Each process can:
  - Request a resource of a given type and block/wait until one resource instance of that type becomes available
  - Use a resource
  - Release a resource

# System Modeling

- Set of resource types:  $R_1, R_2, \dots, R_m$ 
  - There are multiple resource of each type: e.g., 3 NICs, 4 disks
- Set of processes (or threads):  $P_1, P_2, \dots, P_n$
- Each process can:
  - Request a resource of a given type and block/wait until one resource instance of that type becomes available
  - Use a resource
  - Release a resource
- In the previous slide we have two processes,  $P_1$  and  $P_2$  (2 threads), two resource types  $R_1$  (one lock, which corresponds to some resource), and  $R_2$  (another lock, which corresponds to another resource)

# Necessary conditions for deadlock

- That is, what conditions need to be true (of some system) so that deadlock is possible?
  - Aka Coffman conditions
  - 3 necessary conditions ...

# Necessary conditions for deadlock

- That is, what conditions need to be true (of some system) so that deadlock is possible?
  - Aka Coffman conditions
  - 3 necessary conditions ...
- Not the same as causing deadlock!

# Necessary condition 1: Mutual Exclusion

- At least one resource is non-shareable; in other words at most one process at a time can use it
- In our example: the locks are mutually exclusive


## Necessary condition 2: No preemption


- Resources cannot be forcibly removed from processes that are holding them
- In our example: only the thread holding a lock can release it

## Necessary condition 3: Circular wait

- There exists a set  $\{P_0, P_1, \dots, P_p\}$  of waiting processes such that  $(\forall i \in \{0, 1, \dots, p-1\})$   $P_i$  is waiting for a resource held by  $P_{i+1}$  and  $P_p$  is waiting for a resource held by  $P_0$
- In other words, there is a circular chain of processes such that each process holds one or more resources that are being requested by the next process in the chain
- In our example: thread 1 has lock A and needs lock B, and thread 2 has lock B and needs lock A

# Resource-allocation graph

Process: 


Resource: 

Request: 

Allocation: 



# Example Graph




# Cycle and deadlock

- Theorem [Holt]:  
If the resource allocation graph contains no (directed) cycle, then there is no deadlock in the system
  - If cycles do exist then a deadlock is possible

- Theorem [Holt]:  
If the resource allocation graph contains no (directed) cycle, then there is no deadlock in the system
  - If cycles do exist then a deadlock is possible
- If there is only one resource instance (black dot) per resource type then we have a stronger theorem:  
*The existence of a cycle is a necessary and sufficient condition for the existence of a deadlock*


# Cycle and deadlock: our 2-lock example

Clearly, there is a cycle




# Another Example

Can you see the cycle(s)?




# Moving vertices around

Can you see the cycle(s) now?




# Moving vertices around

Can you see the cycle(s) now?



## Example: Cycle and No Deadlock


There is a cycle ...





## Example: Cycle and No Deadlock

There is a cycle ... but there is no deadlock




## Example: Cycle and No Deadlock

When  $P_4$  terminates it will release the instance of  $R_2$  it locked, and that resource will be locked by  $P_3$ .

$P_3$  will then be able to complete.

(Another option is that  $P_2$  completes first.)



# Rule of Thumb

- A cycle in the resource allocation graph
  - Is a necessary condition for a deadlock
  - But not a sufficient condition

# Deadlock Prevention

- Necessary condition 1 – Mutual Exclusion: “At least one resource is non-shareable”
  - If we make this condition not true, then we can't have a deadlock

# Deadlock Prevention

- Necessary condition 1 – Mutual Exclusion: “At least one resource is non-shareable”
  - If we make this condition not true, then we can't have a deadlock
  - In other words: only use shareable resources
    - In our example, if we don't use locks, then we can't have deadlocks :)
  - The problem is that non-shareable resources are useful:
    - A critical section protected by locks, a file open for writing, etc.

# Deadlock Prevention

- Necessary condition 2 – No Preemption: “Resources cannot be forcibly removed”
  - If we make this condition not true, then we can't have a deadlock


# Deadlock Prevention

- Necessary condition 2 – No Preemption: “Resources cannot be forcibly removed”
  - If we make this condition not true, then we can't have a deadlock
  - But how do we even program in an environment in which you acquire resources but then lose them at any time?
  - And would it be safe?


# Deadlock Prevention

- Necessary condition 2 – No Preemption: “Resources cannot be forcibly removed”
  - If we make this condition not true, then we can't have a deadlock
  - But how do we even program in an environment in which you acquire resources but then lose them at any time?
  - And would it be safe?
- In practice, we cannot eliminate mutual exclusion or no-preemption
  - Circular Wait is where it's at.







Possible to create cycle (with one edge)?




Possible to create cycle (with one edge)?




It's quite possible that P2 won't need R2, or, maybe P2 will release R1 before requesting R2, but we don't know if/when ...




Preventing circular wait means avoiding a state where a cycle is an imminent possibility




To predict deadlock, we can ask processes to “claim” all resources they need in advance



Graph with “claim edges”




$P_2$  requests  $R_1$




Convert to allocation edge; no cycle







$P_1$  requests  $R_2$




If we convert to an allocation edge ...




Cycle involving claim edges!




Means that if processes fulfill their claims, we can not avoid deadlock!



$P_1 \rightarrow R_1, P_2 \rightarrow R_2$



$P_1 \rightarrow R_2$  should be blocked by the OS, even if it can be satisfied with available resources




This is a “safe” state ... i.e., no way a process can cause deadlock directly (i.e., without OS alloc)


Idea: if granting an incoming request would create a cycle in a graph with claim edges, deny that request (i.e., block the process)

- approve later when no cycle would occur






P2 releases R1



Now OK to approve  $P_1 \rightarrow R_2$  (unblock  $P_1$ )

Problem: this approach may incorrectly predict imminent deadlock when resources with multiple instances are involved



Requires a more general definition of “safe state”

# Deadlock Detection

- Goal: How can OS detect when there is a deadlock?
- OS should keep track of
  - Current resource allocation (who has what)
  - Current pending requests (who is waiting for what)
- This info is enough to check if there is a current deadlock (see next few slides)

# Detecting Deadlocks

- Suppose there is only one instance of each resource
- Example 1: Is this a deadlock?
  - P1 has R2 and R3, and is requesting R1
  - P2 has R4 and is requesting R3
  - P3 has R1 and is requesting R4
- Example 2: Is this a deadlock?
  - P1 has R2, and is requesting R1 and R3
  - P2 has R4 and is requesting R3
  - P3 has R1 and is requesting R4

# Detecting Deadlocks

- Solution: Build a graph, called Resource Allocation Graph (RAG)
  - There is a node for every process and a node for every resource
  - If process P currently has resource R, then put an edge from R to P
  - If process P is requesting resource R, then put an edge from P to R

# Detecting Deadlocks

- Solution: Build a graph, called Resource Allocation Graph (RAG)
  - There is a node for every process and a node for every resource
  - If process P currently has resource R, then put an edge from R to P
  - If process P is requesting resource R, then put an edge from P to R
- There is a deadlock if and only if RAG has a cycle



# Detecting deadlocks

- How to detect deadlocks when there are multiple instances of resources
- Example:
  - Suppose there are 2 instances of A and 3 of B
  - Process P currently has 1 instance of A, and is requesting 1 instance of A and 3 instances of B
  - Process Q currently has 1 instance of B, and is requesting 1 instance of A and 1 instance of B

# Detecting deadlocks

- Suppose there are  $n$  processes  $P_1, \dots, P_n$  and  $m$  resources  $R_1, \dots, R_m$
- To detect deadlocks, we can maintain the following data structures
  - Current allocation matrix  $C$ :  $C[i,j]$  is the number of instances of resource  $R_j$  currently held by process  $P_i$

# Detecting deadlocks

- Suppose there are  $n$  processes  $P_1, \dots, P_n$  and  $m$  resources  $R_1, \dots, R_m$
- To detect deadlocks, we can maintain the following data structures
  - Current allocation matrix  $C$ :  $C[i,j]$  is the number of instances of resource  $R_j$  currently held by process  $P_i$
  - Current request matrix  $R$ :  $R[i,j]$  is the number of instances of resource  $R_j$  currently being requested by process  $P_i$

# Detecting deadlocks

- Suppose there are  $n$  processes  $P_1, \dots, P_n$  and  $m$  resources  $R_1, \dots, R_m$
- To detect deadlocks, we can maintain the following data structures
  - Current allocation matrix  $C$ :  $C[i,j]$  is the number of instances of resource  $R_j$  currently held by process  $P_i$
  - Current request matrix  $R$ :  $R[i,j]$  is the number of instances of resource  $R_j$  currently being requested by process  $P_i$
  - Availability vector  $A$ :  $A[j]$  is the number of instances of resources  $R_j$  currently free.

# Deadlock detection

- Goal of the detection algorithm is to check if there is any sequence in which all current requests can be met
  - Note: If a process  $P_i$ 's request can be met, then  $P_i$  can potentially run to completion, and release all the resources it currently holds. So for detection purpose,  $P_i$ 's current allocation can be added to  $A$

# Example

$L = \{\}$  /\* List of processes that can be unblocked \*/

Allocation Matrix C

		R1	R2	R3
-----				
P1		1	1	1
P2		2	1	2
P3		1	1	0
P4		1	1	1

Request Matrix R

		R1	R2	R3
-----				
P1		3	2	1
P2		2	2	1
P3		0	0	1
P4		1	1	1

Satisfiable request

$A = (0, 0, 1)$  /\* available resources \*/

Request by process  $i$  can be satisfied if the row  $R[i]$  is smaller than or equal to the vector  $A$

# After first iteration

$$L = \{P3\}$$

Allocation Matrix

| R1 R2 R3

-----  
P1 | 1 1 1

P2 | 2 1 2

P3 |

P4 | 1 1 1

$$A = (1, 1, 1)$$

Request Matrix

| R1 R2 R3

-----  
P1 | 3 2 1

P2 | 2 2 1

P3 |

P4 | 1 1 1

Satisfiable request

Note: P3's allocation has been added to A

## After second iteration

$$L = \{P3, P4\}$$

Allocation Matrix			
	R1	R2	R3
P1	1	1	1
P2	2	1	2
P3			
P4			

Request Matrix			
	R1	R2	R3
P1	3	2	1
P2	2	2	1
P3			
P4			

Satisfiable request

$$A = (2, 2, 2).$$



# Algorithm

```
L = EmptyList; /* processes not deadlocked */
repeat
    s = length(L);
    for (i=1; i<=n; i++){
        if (!member(i,L) && R[i] <= A) {
            /* request of process i can be met */

            A = A + C[i];
            /* reclaim resources held by process i */

            insert(i,L);
        }
    }
until (s == length(L));
/* if L does not change, then done */

if (s<n) printf("Deadlock exists");
```

Note: Running time of this algorithm is  $O(n^2m)$ , where  $m$ : length of a row

# Last time

- Each process provides OS with information about its requests and releases for resources
- OS determines if there is any sequence in which all current requests can be met, avoiding deadlocks

- Each process provides OS with information about its requests and releases for resources
- OS determines if there is any sequence in which all current requests can be met, avoiding deadlocks
- Assumption: processes can complete with current allocation + all current requests

- Each process provides OS with information about its requests and releases for resources
- OS determines if there is any sequence in which all current requests can be met, avoiding deadlocks
- Assumption: processes can complete with current allocation + all current requests
- i.e., no future requests
  - Unrealistic!

# Deadlock avoidance

- Knowing all future requests is difficult
- Estimating a maximum demand for resources for each process is easier to produce

# Deadlock avoidance

- Knowing all future requests is difficult
- Estimating a maximum demand for resources for each process is easier to produce
- simple strategy: specify a maximum claim
- A resource allocation state is defined by
  - no. of available resources
  - no. of allocated resources to each process

# Deadlock avoidance

- Knowing all future requests is difficult
- Estimating a maximum demand for resources for each process is easier to produce
- simple strategy: specify a maximum claim
- A resource allocation state is defined by
  - no. of available resources
  - no. of allocated resources to each process
  - maximum demands by each process

# Deadlock avoidance


- A state is safe if there exists a safe sequence of processes  $\langle P_1, \dots, P_n \rangle$  for the current resource allocation state
  - A sequence of processes is safe if for each  $P_i$  in the sequence, the resource requests that  $P_i$  can still make can be satisfied by:
    - currently available resources + all resources held by all previous processes in the sequence  $P_j, j < i$
  - If resources needed by  $P_i$  are not available,  $P_i$  waits for all  $P_j$  to release their resources



# Deadlock avoidance


- Intuition for a safe state: given that the system is in a certain state, we want to find at least one “way out of trouble”
- i.e. find a sequence of processes that, even when they demand their maximum resources, won't deadlock the system

# Deadlock avoidance



- A deadlocked state is unsafe
- An unsafe state is not necessarily deadlocked

# Deadlock avoidance



- A deadlocked state is unsafe
- An unsafe state is not necessarily deadlocked
- A system may transition from a safe to an unsafe state if a request for resources is granted

# Deadlock avoidance

- Example 1:
  - 12 instances of a resource
  - At time  $t_0$ ,  $P_0$  holds 5,  $P_1$  holds 2,  $P_2$  holds 2
  - Available = 3 free instances

processes	max needs	allocated
P0	10	5
P1	4	2
P2	9	2

- Example 1 (cont)
  - Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
    - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource

- Example 1 (cont)

- Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
  - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
  - Then  $P_1$  then releases all of its held resources, so that there are 5 free

- Example 1 (cont)

- Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
  - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
  - Then  $P_1$  then releases all of its held resources, so that there are 5 free
  - Next, suppose  $P_0$  requests its maximum (currently has 5, so needs 5 more) and holds 10, so that there are 0 free

- Example 1 (cont)

- Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
  - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
  - Then  $P_1$  then releases all of its held resources, so that there are 5 free
  - Next, suppose  $P_0$  requests its maximum (currently has 5, so needs 5 more) and holds 10, so that there are 0 free
  - Then  $P_0$  releases all of its held resources, so that there are 10 free



- Example 1 (cont)

- Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
  - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
  - Then  $P_1$  then releases all of its held resources, so that there are 5 free
  - Next, suppose  $P_0$  requests its maximum (currently has 5, so needs 5 more) and holds 10, so that there are 0 free
  - Then  $P_0$  releases all of its held resources, so that there are 10 free
  - Next  $P_2$  requests its max of 9, leaving 3 free and then releases them all

- Example 1 (cont)

- Claim: the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe.
  - $P_1$  requests its maximum (currently has 2, so needs 2 more) and holds 4, then there is only 1 free resource
  - Then  $P_1$  then releases all of its held resources, so that there are 5 free
  - Next, suppose  $P_0$  requests its maximum (currently has 5, so needs 5 more) and holds 10, so that there are 0 free
  - Then  $P_0$  releases all of its held resources, so that there are 10 free
  - Next  $P_2$  requests its max of 9, leaving 3 free and then releases them all
- Thus the sequence  $\langle P_1, P_0, P_2 \rangle$  is safe

# Deadlock avoidance

- Example 2:

- at time  $t_1$ , process  $P_2$  requests and is given 1 more instance of the resource, then

processes	max needs	allocated
P0	10	5
P1	4	2
P2	9	3

- Available = 2 free instances
- Is the system in a safe state?

- Example 2 (cont)
  - $P_1$  can request its maximum (currently holds 2, and needs 2 more) of 4, so that there are 0 free

- Example 2 (cont)
  - $P_1$  can request its maximum (currently holds 2, and needs 2 more) of 4, so that there are 0 free
  - $P_1$  releases all its held resources, so that available = 4 free

- Example 2 (cont)

- $P_1$  can request its maximum (currently holds 2, and needs 2 more) of 4, so that there are 0 free
- $P_1$  releases all its held resources, so that available = 4 free
- Neither  $P_0$  nor  $P_2$  can request their maximum resources ( $P_0$  needs 5,  $P_2$  needs 6, and there are only 4 free)
  - Both would have to wait, so there could be deadlock

- Example 2 (cont)

- $P_1$  can request its maximum (currently holds 2, and needs 2 more) of 4, so that there are 0 free
- $P_1$  releases all its held resources, so that available = 4 free
- Neither  $P_0$  nor  $P_2$  can request their maximum resources ( $P_0$  needs 5,  $P_2$  needs 6, and there are only 4 free)
  - Both would have to wait, so there could be deadlock
- The system is deemed unsafe
  - The mistake was granting  $P_2$  an extra resource at time  $t_1$
  - Forcing  $P_2$  to wait for  $P_0$  or  $P_1$  to release their resources would have avoided potential deadlock

# Deadlock avoidance

- Policy: before granting a request, at each step, perform a worst-case analysis - is there a safe sequence, a way out?
  - If so, grant request.
  - If not, delay requestor, and wait for more resources to be freed.
- Banker's Algorithm takes this approach



- Banker's Algorithm:

- when there is a request, the system determines whether allocating resources for the request leaves the system in a safe state that avoids deadlock
  - if no, then wait for another process to release resources
- each process declares its maximum demands
  - must be less than total resources in system
- works for multiple instances of resources, unlike resource allocation graph

# Deadlock avoidance

- Define quantities:
  - Available resources in a vector or list  $\text{Available}[j]$ ,  $j=1,\dots,m$  resource types
    - $\text{Available}[j] = k$  means that there are  $k$  instances of resource  $R_j$  available
  - Maximum demands in a matrix or table  $\text{Max}[i,j]$ , where  $i=1,\dots,n$  processes, and  $j=1,\dots,m$  resource types
    - $\text{Max}[i,j] = k$  means that process  $i$ 's maximum demands are for  $k$  instances of resource  $R_j$
  - Allocated resources in a matrix  $\text{Alloc}[i,j]$ 
    - $\text{Alloc}[i,j] = k$  means that process  $i$  is currently allocated  $k$  instances of resource  $R_j$
  - Needed resources in a matrix  $\text{Need}[i,j] = \text{Max}[i,j] - \text{Alloc}[i,j]$

# Deadlock avoidance

- An example of the  $\text{Alloc}[i,j]$  matrix:

		Resources					
		Column j					
Processes	Row i		1				
			7				
			12				
		8	0	2	17	0	1
			0				
			4				
			0				
			1				

←  $\text{Alloc}_i$  is shorthand for row  $i$  of matrix  $\text{Alloc}[i,j]$ , i.e. the resources allocated to process  $i$

# Deadlock avoidance

- Some terminology:

- let  $X$  and  $Y$  be two vectors. Then we say  $X \leq Y$  if and only if  $X[i] \leq Y[i]$  for all  $i$ .
- Example:

$$V1 = \begin{bmatrix} 1 \\ 7 \\ 3 \\ 2 \end{bmatrix}$$

$$V2 = \begin{bmatrix} 0 \\ 3 \\ 2 \\ 1 \end{bmatrix}$$

$$V3 = \begin{bmatrix} 0 \\ 10 \\ 2 \\ 1 \end{bmatrix}$$

then  $V2 \leq V1$ , but  
 $V3 \not\leq V1$ , i.e.  $V3$  is  
not less than or equal  
to  $V1$

# Deadlock avoidance

- Banker's (Safety) Algorithm: find a safe sequence, i.e. is the system in a safe state?
  - ① Let Work and Finish be vectors length m and n respectively.  
Initialize Work = Available, and Finish[i]=false for i=0, ..., n-1
  - ② Find a process i such that both
    - Finish[i]==false, and
    - $Need_i \leq Work$If no such i exists, go to step 4.
  - ③ Work = Work+ Alloc;  
Finish[i] = true  
Go to step 2.
  - ④ If Finish[i]==true for all i, then the system is in a safe state

# Deadlock avoidance

- Example 3:

- 3 resources (A,B,C) with total instances available (10,5,7)
- 5 processes
- At time  $t_0$ , the allocated resources  $Alloc[i,j]$ , Max needs  $Max[i,j]$ , and Available resources  $Avail[j]$ , are:

	Alloc[i,j]				Max[i,j]				Avail[j]				Need[i,j]		
	A	B	C		A	B	C		A	B	C		A	B	C
P0	0	1	0		7	5	3						7	4	3
P1	2	0	0		3	2	2						1	2	2
P2	3	0	2		9	0	2						6	0	0
P3	2	1	1		2	2	2						0	1	1
P4	0	0	2		4	3	3						4	3	1

Avail[j]

A	B	C
3	3	2

where  $Need[i,j]$  is  
computed given  
 $Alloc[i,j]$  and  $Max[i,j]$

- Example 3 (cont):
  - Execute Banker's Algorithm – is the system in a safe state?
  - Yes, the sequence  $\langle P1, P3, P4, P2, P0 \rangle$  is safe
  - In-class Exercise